

Minnie Nguyen

☎ +1-438-338-7557 | ✉ minniexnguyen@gmail.com | 🌐 nminnie | 🐱 nminnie | 🌐 nminnie.github.io

AI Engineer (MSc with Distinction, University of Edinburgh) with 5+ years building NLP/LLM and ML systems in production, spanning fine-tuning, evaluation, and deployment. Skilled in translating recent research into reliable, scalable AI solutions.

EDUCATION

University of Edinburgh

Edinburgh, UK

Master of Science in Artificial Intelligence (Distinction)

Sep 2024 – Aug 2025

- **Thesis:** Systematic framework to quantify modality reliance in multilingual Vision-Language Models [paper]
- **Research Project:** Data-centric object detection pipeline for marine debris using YOLO/OpenCV [paper]
- **Research Review:** Leveraging auxiliary languages for low-resource Neural Machine Translation [paper]

McGill University

Montreal, Canada

Bachelor of Science in Mathematics and Computer Science, GPA: 3.73/4.0

Sep 2015 – May 2019

- **Awards:** Tomlinson Award for Mentoring (2017, 2018, 2019), Major Entrance Scholarship (2015)

EXPERIENCE

Got It AI

Hanoi, Vietnam (Remote US)

Machine Learning Research Engineer

Feb 2023 – May 2024

- Developed AI hallucination detection guardrails using multi-task training with chain-of-thought reasoning based on Google's *Distilling Step-by-Step* method, improving recall by 21% to ensure reliable LLM outputs for end users.
- Fine-tuned LLMs (Mistral, LLaMA, Flan-T5) using LoRA-based knowledge distillation and SFT to build specialized conversational AI assistants for global enterprise clients (USA, India, Hong Kong).
- Deployed fine-tuned LLMs on GCP using Docker, FastAPI, and vLLM with 8-bit quantization, optimizing GPU memory usage while preserving task accuracy.
- Built an active learning pipeline to retrain models on user interaction data, reducing error by up to 5% per iteration and driving continuous model improvement.
- Benchmarked state-of-the-art LLMs (OpenAI, Gemini, Mixtral) on curated question-answering datasets from client knowledge bases, informing model selection and deployment decisions.

Intact Insurance, Data Lab

Montreal, Canada

Data Scientist

Sep 2020 – Dec 2022

Data Scientist Intern

May 2020 – Aug 2020

- Fine-tuned BERT-based models for text classification and NER on unstructured English and French claim documents, achieving 90%+ accuracy to automate insurance claim analysis at scale.
- Led domain adaptation of language models by further pretraining on proprietary insurance documents, improving F1-scores in 3+ downstream NLP projects by 5-15%.
- Implemented a stacking ensemble pipeline (LSTM, XGBoost, LightGBM) adapted from UC Berkeley's *Super Learner* method with advanced temporal feature engineering to forecast claim payment amounts.

Dynamicly Inc.

Montreal, Canada

Data Scientist

Oct 2019 – Apr 2020

Data Scientist Intern

Jun 2019 – Sep 2019

- Designed and built a NLU system using text embeddings and spaCy/NLTK to power virtual assistants, outperforming Microsoft LUIS by 29% in intent classification and entity extraction accuracy.
- Collaborated with engineering and product teams to integrate models into production workflows as a REST API serving multiple virtual assistant applications.
- Created analytics dashboards in Power BI to monitor model performance and extract user interaction insights, enabling data-driven improvements to the NLU pipeline.

SKILLS

Programming: Python, Bash, SQL, Java, JavaScript, HTML/CSS

Libraries: PyTorch, HuggingFace Transformers, LoRA/QLoRA, Scikit-Learn, Pandas, OpenCV, Ultralytics

Tools: Git, GCP, CUDA, Docker, FastAPI, vLLM, W&B, MLflow, Jupyter

Expertise: LLM Fine-tuning, Evaluation, Multilingual NLP, Computer Vision, Multimodal VLMs

Languages: English (fluent), Vietnamese (fluent), French (intermediate)